

Exercise 3. Let W be a walk from x and y . Prove that W contains a path P from x and y . Furthermore, if the edge weights are nonnegative, prove that P 's total cost is at most W 's total cost.

- We can reduce the walk W to be the path P
- So, with nonnegative edge weights, P 's total cost is at most W , both traversing the same set of edges and nodes.

Let W be a walk from x to y . We consider two cases:

[Case 1: W does not repeat any vertices] Since W does not repeat any vertices for this case, then it is already a path P .

[Case 2: W repeats vertices] Suppose some vertex v appears at least twice in W . Consider the first two occurrences of v . The portion of W between these occurrences forms a cycle. Removing this cycle yields another walk W' from x to y .

Repeat this process whenever a vertex appears more than once. Each removal shortens the walk, so the process terminates. The resulting walk has no repeated vertices and is therefore a path P from x to y .

Problem 4. Single Source Shortest Path. Given a connected weighted digraph, and a node x , what is the length of the shortest path from x to y , for every node y ?

Problem 5. Single Destination Shortest Path. Given a connected weighted digraph, and a node x , what is the length of the shortest path from y to x , for every node y ?

Problem 6. All-Pair Shortest Path. Given a connected weighted digraph, what is the length of the shortest path from x to y , for every node x and every node y ?

Exercise 7. Reduce Problem 4 to Problem 5 and vice versa. That is, find an algorithm for Problem 4 assuming you have a magic oracle that can solve instances of Problem 5, and vice versa.

We can solve SSSP via an SDSP oracle by reversing the graph, then computing the distance to x , which will be equal to the distance to x .

To solve SDSP via an SSSP oracle, we do the same: reverse the graph, then compute the distance to every node y .

Exercise 8. Reduce Problem 6 to Problem 4 have an algorithm that solves Problem 4 with only a $\Theta(n)$ overhead. That is, assuming you an algorithm that solve Problem 4 in $\mathcal{O}(f(n))$ time, find an algorithm that solves Problem 6 in $\mathcal{O}(n \cdot f(n))$ time.

We run the SSSP algorithm once per vertex. The i -th run computes the distances from vertex i to all others, so running it for all n nodes would solve Problem 6 since it would yield the distances for all pairs. Since we perform n runs, the algorithm would run in $O(n * f(n))$ time.

Why do we drop the k variable in Floyd-Warshall?

The invariant for Floyd-Warshall's algorithm is: "At the start of iteration k , $d[i][j]$ equals the length of the shortest path from i to j whose immediate vertices are $\{0, 1, \dots, k - 1\}$. Then $\{0, 1, \dots, k\}$ after iteration k finishes."

We can update the matrix in place because during iteration k , we only introduce vertex k as a new allowed intermediate. Since distances only decrease and don't use vertices beyond k , the DP invariant remains valid.

How can we detect cycles in Floyd-Warshall?

We can detect cycles in Floyd-Warshall by checking whether any diagonal entry is negative. The length of the shortest path from a vertex to itself is 0, so if $d[i][i] < 0$, then a negative cycle is reachable from the vertex i .