

### **done with 1 and 2 in submissions, starting from 3**

3. In the handouts, it was shown that the height of an AVL tree with  $n$  nodes has height  $\mathcal{O}(\log n)$ , by first proving that  $N(h) \geq \alpha\beta^h + \gamma$  for some constants  $\alpha > 0$ ,  $\beta > 1$  and  $\gamma$ , which directly implies that  $h \leq \frac{\lg n}{\lg \beta} + \mathcal{O}(1)$ , and thus  $h = \mathcal{O}(\log n)$ .

The actual value of  $\beta$  we got from the handouts is  $\beta = \sqrt{2}$ , which gives us the tighter bound  $h \leq \frac{\lg n}{\lg \sqrt{2}} + \mathcal{O}(1) = 2 \lg n + \mathcal{O}(1)$ .

The goal of this task is to prove an even tighter bound, by proving that we can choose  $\beta$  to be  $\varphi$ .<sup>1</sup> This gives us a height bound of  $\approx 1.44 \lg n + \mathcal{O}(1)$ .

#### **3.1. Let $f$ be defined recursively as $f(0) = 1$ , $f(1) = 2$ and for $n \geq 2$ ,**

$$f(n) = f(n-1) + f(n-2) + 1.$$

**Prove that  $N(h) \geq f(h)$  for all  $h \geq 0$ .**

We prove by strong induction on  $h$ .

[Base Cases]

$$h = 0: N(0) = 1 = f(0)$$

$$h = 1: N(1) = 2 = f(1)$$

Thus,  $N(h) \geq f(h)$  holds.

[Inductive Step] Assume for all  $k < h$  that  $N(k) \geq f(k)$ .

For  $h \geq 2$ , an AVL tree of height  $h$  with minimal nodes must have one child of height  $h-1$  and the other of height at least  $h-2$ . Hence:

$$N(h) \geq N(h-1) + N(h-2) + 1$$

By the inductive hypothesis,  $N(h) \geq f(h-1) + f(h-2) + 1$ .

But by definition of  $f$ :  $f(h-1)+f(h-2)+1=f(h)$ .

Therefore,  $N(h) \geq f(h)$ , completing the induction and the proof.

**3.2. Show that the equation  $x^2 = x + 1$  has exactly two solutions, both of which are real numbers, and find the largest one. Call this constant b.**

Suppose  $r$  is a solution, so that  $r^2 = r + 1$ . Then we have:

$$\begin{aligned}r^2 - r &= 1 \\4r^2 - 4r &= 4 \\4r^2 - 4r + 1 &= 5 \\(2r - 1)^2 &= 5 \\2r - 1 &= \{\sqrt{5}, -\sqrt{5}\} \\r &\in \{(1+\sqrt{5})/2, (1-\sqrt{5})/2\}\end{aligned}$$

Thus, the only possible solutions are  $(1+\sqrt{5})/2, (1-\sqrt{5})/2$ , we can verify that these solutions work. We pick the largest of the two:  $1+\sqrt{5}/2 = \phi$ .

**3.3. Show that there exists a constant  $a > 0$  such that  $f(n) \geq ab^n$  for all  $n \geq 0$ .**

We have  $b = \phi$ , the positive root of  $x^2 = x + 1$ . Choose  $a = 1$ . We prove by strong induction that  $f(n) \geq b^n$  for all  $n \geq 0$ .

[Base Cases]

$$n = 0: f(0) = 1 \geq b^0 = 1$$

$$n = 1: f(1) = 2 \geq b^1 \approx 1.618$$

Thus, the base cases hold for  $n = 0, 1$ .

[Inductive Step] Assume for all  $k < n$  that  $f(k) \geq b^k$ . For  $n \geq 2$ , by definition of  $f$ :

$$f(n) = f(n-1) + f(n-2) + 1.$$

By the inductive hypothesis:  $f(n) \geq b^{n-1} + b^{n-2} + 1$ .

since  $b^2 = b + 1$ , we have  $b^{n-1} + b^{n-2} = b^{n-2}(b+1) = b^{n-2} b^2 = b^n$

Thus,  $f(n) \geq b^n + 1 > b^n$

Therefore,  $f(n) \geq b^n$  for all  $n \geq 0$ , proving the claim with  $a = 1$ .

**3.4. Show that  $N(h) \geq ab^h$  for all  $h \geq 0$  (for the  $a$  chosen in 3.3)**

Combining 3.1 and 3.3, we get  $N(h) \geq f(h) \geq ab^h$  for every  $h \geq 0$ .

### 3.5. Conclude that an AVL tree with $n$ nodes has height $\leq c \lg n + O(1)$ where $c = 1/\lg \phi$ .

Let  $T$  be a tree with  $n$  nodes, and suppose its height is  $h$ . Note that  $n \geq N(h)$  by definition of  $N$ . Using 3.4, we have

$$n \geq ab^h = \phi^h$$

Taking logarithms, we get  $\lg n \geq h \lg \phi$ , implying

$$h \leq 1/\lg \phi * \lg n,$$

which is what we wanted to prove.

4. A **red-black tree** is another kind of self-balancing tree, different from an AVL tree. It maintains balancing by coloring each node either *red* or *black*, and maintaining some invariants. To maintain these invariants, tree rotations are used, similar to AVL trees.

Its height is also  $\mathcal{O}(\log n)$ , though the constant behind the “ $\mathcal{O}$ ” here is generally larger than that in the case of the AVL tree, so red-black trees are often taller (though just by a constant). However, it has a few notable “selling points” as well:

- The only extra data needed to be stored in each node is its color. Since there are only two colors, only one bit of extra data per node is needed.<sup>2</sup>
- Although insertion and deletion as a whole still take  $\mathcal{O}(\log n)$  time, only  $\mathcal{O}(1)$  tree rotations are needed to rebalance the tree!
- It is related to B-trees, which are yet another kind of self-balancing tree, particularly suited for databases and file systems.

The following invariants are maintained after every operation in a red-black tree:

- 4.1. A red node does not have a red child.
- 4.2. If a left or right child of a node is missing, we say that its left or right child is a null node.
- 4.3. A null node has no children and is colored black.
- 4.4. All paths from the root to a null node have the same number of black nodes.

### 4. Prove that any tree with $n$ nodes satisfying the red-black tree properties above has height $\leq 2 \lg n + O(1)$ .

Let  $c$  be the black-height of the tree, the number of black nodes on any path from root to null.

By induction, any red-black tree with black-height  $c$  has at least  $2^c - 1$  nodes  $\geq 2(2^{c-1} - 1) + 1 = 2^c - 1$ .

If root is red, then both children are black with black-height  $c$ , with even more nodes. Thus,  $n \geq 2^c - 1$ . On any root-to-null path, no two reds are consecutive, so at least half the nodes are black (rounded up).

Therefore,  $h \leq 2c$ .

Combining  $h \leq 2c \leq 2\log_2(n+1) = 2\lg n + O(1)$ . Hence, height  $\leq 2\lg n + O(1)$ .

5. An instance of the **dictionary** abstract data type consists of a collection of key-value pairs, where the keys are distinct. For simplicity, the keys and values in our dictionary will be strings. It supports the following operations:
  - 5.1. Create an empty dictionary.
  - 5.2. Find the size (number of key-value pairs) in the dictionary.
  - 5.3. Insert a new key-value pair to the dictionary (unless a pair with that key already exists, in which case, do nothing).
  - 5.4. Delete a key-value pair with a given key (unless there's no such pair, in which case, do nothing).
  - 5.5. Find the corresponding value of a given key.

## **5. Explain how to use AVL trees to implement a dictionary such that the first two operations take $O(1)$ time, and the last three operations take $O(\log n)$ time.**

For this data structure, we store key-value pairs in an AVL tree, ordered by key. Each node contains: key, value, left, right, height. We maintain a separate size attribute, tracking the total number of nodes.

- 5.1. To create an empty dictionary, we initialize  $\text{root} = \text{null}$  and  $\text{size} = 0$ . This takes  $O(1)$ .
- 5.2. Provided we've maintained the size correctly throughout the operations, all we need to do is return the size attribute. This takes  $O(1)$ .
- 5.3. To insert a key-value pair into the AVL tree:
  - Perform standard AVL insert using key for ordering.
  - If the key exists, update its value, return False.
  - If the key does not exist, create a new node, increment size by 1, and return True.
  - After insertion, perform rebalancing along the path.
  - This takes  $O(\log n)$ .
- 5.4. To delete a key-value pair with a given key:
  - Perform standard AVL delete using key.
  - If key found: remove node, decrement size by 1, return True.
  - If key not found: return False
  - After deletion, perform rebalancing along the path.
  - This takes  $O(\log n)$ .

5.5. To find a value given a key:

- We perform standard BST search using key.
- If found: return associated value.
- If not, return null or identity value.
- This takes  $O(\log n)$ .

AVL property ensures height  $O(\log n)$ , so all tree operations run in  $O(\log n)$ . The size variable is updated on every insertion/deletion, so it is maintained correctly.

6. Consider an abstract data type that represents a collection of integers and that supports the following operations:

- 6.1. Given an integer  $x$ , insert  $x$  in the structure. Note that the same integer can be inserted multiple times.
- 6.2. Given an integer  $x$ , remove  $x$  from the structure. If there are multiple occurrences of  $x$ , remove only one occurrence. If there are none, do nothing.
- 6.3. Given an integer  $x$ , determine how many elements of the structure are strictly less than  $x$ .

When an instance is created, it is initially empty.

Describe a data structure implementing the above, with each operation running in  $\mathcal{O}(\log n)$  worst case time.

## 6. Describe a data structure implementing the above, with each operation running in $O(\log n)$ worst case time.

We implement an AVL tree, augmented as follows:

- Each node stores: value, count (multiplicity), left, right, height, size (total nodes in subtree)
- Multiple occurrences of same value stored in a single node with count  $> 1$ .

For node  $x$ :

$$\begin{aligned}\text{size}(x) &= \text{size(left)} + \text{size(right)} + \text{count}(x) \\ \text{height}(x) &= \max(\text{height(left)}, \text{height(right)}) + 1\end{aligned}$$

We update the parameters above during rotations/rebalancing in  $O(1)$ .

6.1. Inserting  $x$ :

- Search for node with value  $x$
- If found: increment count by 1
- If not found: create new node with count = 1
- Update size along path, rebalance if needed.
- This takes  $O(\log n)$ .

### 6.2. Removing one occurrence of x:

- Search for node with value x
- If not found: do nothing
- If found and count > 1: decrement count by 1
- If found and count = 1 delete node, using the standard AVL deletion.
- Update size along path, rebalance if needed.
- This takes  $O(\log n)$ .

### 6.3. Count elements strictly lesser than x:

- We use the function below:

```

```
function count_less(node, x):
    if node is null: return 0
    if x <= node.value:
        return count_less(node.left, x)
    else: // x > node.value
        return size(node.left) + node.count + count_less(node.right, x)
````
```

- Running it from the root, and traversing at most one path.
- This takes  $O(\log n)$ .

We now discuss the correctness of this implementation. AVL property guarantees height  $O(\log n)$ . The size augmentation enables  $O(\log n)$  order statistics. All operations update size/height correctly during rotations.