

Exercise 1. Prove that any connected graph with $n \geq 1$ node has $\geq n - 1$ edges. (Hint: Adding an edge can only decrease the number of connected components by 1.)

We prove by induction on the number of edges.

[Base Case] A connected graph with $n = 1$ has 0 edges, and $0 \geq 1 - 1 = 0$, so the statement holds.

[Inductive Case] Assume any connected graph with n nodes and m edges satisfies $m \geq n - 1$. We take a connected graph G with n nodes and $m + 1$ edges. Remove any edge e from G .

[Case 1: G remains connected] Then, by the induction hypothesis, $m \geq n - 1$, so $m + 1 \geq n$.

[Case 2: G becomes disconnected] Then, $G - e$ now has two connected components. Let them have $n_1 + n_2 = n$ nodes.

By the inductive hypothesis, the total edges in $G - e$ is at least $(n_1 - 1) + (n_2 - 1) = n - 2$. Then, G has at least $(n - 2) + 1 = n - 1$ edges.

Thus, any connected graph with n nodes has at least $n - 1$ edges.

Exercise 2. Prove that any acyclic graph with $n \geq 1$ nodes has $\leq n - 1$ edges. (Hint: Adding a bridge decreases the number of connected components by 1.)

We prove by induction on the number of nodes n .

[Base Case] A graph with $n = 1$ has at most 0 edges, and $0 \leq n - 1 = 1 - 1 = 0$, so the statement holds.

[Inductive Case] Assume any acyclic graph with fewer than n nodes has at most $(n - 1) - 1 = n - 2$ edges. Take an acyclic graph G with n nodes and m edges.

since G is acyclic, it has at least one leaf (a node of degree 1). Remove that leaf and its incident edge. The remaining graph has $n - 1$ nodes, remains acyclic, and has $m - 1$ edges.

By the inductive hypothesis, $m - 1 \leq (n - 1) \leq n - 2$

Thus, $m \leq n - 1$.

Hence, any acyclic graph with n nodes has at most $n - 1$ edges.

Exercise 3. Prove that a tree with $n \geq 1$ nodes has exactly $n - 1$ edges.

A tree is defined as a connected and acyclic graph. By Exercise 1, any connected graph with n nodes has at least $n - 1$ edges. By Exercise 2, any acyclic graph with n nodes, has at most $n - 1$ edges.

since a tree is both connected and acyclic, it must satisfy both inequalities simultaneously:

$$n - 1 \leq (\text{number of edges}) \leq n - 1$$

Therefore, a tree with n nodes has exactly $n - 1$ edges.

Exercise 4. Prove that a graph is a tree if and only if it satisfies at least two of the following:

1. It is connected.
2. It is acyclic.
3. It has $n - 1$ edges.

Note that a tree satisfies all three.

[\Rightarrow] “If it satisfies two of the three properties, then the graph is a tree.”

We consider the following cases:

[Case 1: Connected and Acyclic] Then by definition, the graph is a tree, and by Exercises 1 and 2, it also has $n - 1$ edges.

[Case 2: Connected and $n - 1$ edges] suppose for contradiction, that the graph has a cycle. Removing an edge from the cycle keeps it connected, but leaves $n - 2$ edges, contradicting Exercise 1, which states that for a graph to be connected, it has to have at least $n - 1$ edges. Hence, it's also acyclic and is also a tree.

[Case 3: Acyclic and $n - 1$ edges] suppose the graph is disconnected with $k \geq 2$ components, each a tree. Let component i have n_i nodes, so the sum of $n_i = n$ and total edges, are the sum of $(n_i - 1) = n - k$. But given that the graph has $n - 1$ edges,

$$n - k = n - 1 \Rightarrow k - 1$$

Therefore, it is actually just one component, the graph is connected.

Thus, any of the two imply the third, and a tree, which is connected and acyclic, satisfies all three.

[\Leftarrow] “If the graph is a tree, then it satisfies at least two of the three properties.” By definition, a tree is a connected, acyclic graph with n nodes and $n - 1$ edges.

The following exercises show that there is a unique path from one node to another node in a tree.

Exercise V. Prove that if there is a walk from x to y , then there is a path from x to y using only a subsequence of the edges of that walk.

Let $W = (x = v_0, v_1, \dots, v_k = y)$ be a walk from x to y . We consider the following cases, regarding the occurrences of vertices:

[Case 1: No vertex repeated in W] If no vertex is repeated in W , then W is already a path.

[Case 2: Some vertex v_i is repeated] If some vertex $v_i = v_j$ with $i < j$, we remove the cycle v_i, v_{i+1}, \dots, v_j from the walk to get a shorter walk from x to y , and repeat until no vertex is repeated. The result is a path from x to y whose edges are a subsequence of the edges in W .

Therefore, if there is a walk from x to y , then there is a path from x to y .

Exercise 6. Prove that if there are two distinct paths from x to y , then the graph has a cycle. (Hint: Find a nontrivial circuit first).

Let P_1 and P_2 be two distinct paths from x to y . Follow P_1 from x to y , then follow P_2 backward from y to x . This gives a closed walk, a circuit.

If the two paths share no common vertices except x and y , this closed walk is a cycle. If they share other vertices, the closed walk contains a repeated vertex, so it can be reduced to a cycle.

Therefore, the graph contains a cycle.

Exercise 7. Prove that there is at most one path between any two nodes in an acyclic graph. (Hint: Use Exercise 6).

Suppose, for contradiction, that there are two distinct paths between two nodes a and b in an acyclic graph.

By Exercise 6, the existence of two distinct paths from a to b implies the graph contains a cycle, which contradicts the assumption that the graph is acyclic.

Therefore, there is at most one path between any two nodes in an acyclic graph.

E%erci5e 8. Prove that there i5 a unique path between any two node5 in a tree.

Since a tree i5 acyclic, then by E%erci5e 7, there i5 at mo5t one path (unique) between any two node5 in a tree.

E%erci5e 9. Let u be an ance5tor of v . Prove that $\text{di5t}(u, v) = v.\text{depth} - u.\text{depth}$.

Since u i5 an ance5tor of v , the unique path from u to v goe5 down the tree along the parent-child edge5 from u to v .

By definition, $v.\text{depth}$ i5 the number of edge5 from the root to v , and $u.\text{depth}$ i5 the number of edge5 from the root to u .

The number of edge5 between them, i5 the number of edge5 from u to v , and i5 indeed, the difference:

$$\text{di5t}(u, v) = v.\text{depth} - u.\text{depth}$$

This works because each integer c can be written in binary, so the code above eventually removes each bit of c from MSB to LSB. Similarly, we can optimize lca by again going through these pointers in decreasing size while we can:

```
def lca(u, v):
    u = climb(u, u.depth - v.depth)
    v = climb(v, v.depth - u.depth)

    for k from h to 0:
        if u.anc[k] != v.anc[k]:
            u = u.anc[k]
            v = v.anc[k]

        # still need to go up once
        if u != v:
            u = u.parent
            v = v.parent

    return u
```

Exercise 12. Explain why the last step of using the parent pointer once is necessary.

E%erci5e 12. E%plain why the la5t 5tep of u5ing the parent pointer once i5 nece55ary.

In the LCA algorithm using binary lifting, after the ‘for’ loop, u and v are the deepest nodes whose ancestors are not common ancestors of the original u and v . At this point, $u.\text{anc}[0]$ and $v.\text{anc}[0]$, their parents are the same node, the LCA.

Thus, one final step up the parent pointer is needed to reach the LCA from both u and v .



Note that “Euler tour” is not quite the same as an “Eulerian circuit”, because each edge is traversed twice—once upward, and once downward.⁴ Performing an Euler tour is straightforward, and afterward, we get a list of nodes in the Euler tour.

For example, for the tree above, the Euler tour is:

$$E = [0, 1, 0, 2, 0, 3, 4, 3, 5, 3, 0].$$

As it turns out, the LCA of u and v can be obtained as follows:

1. Let i be the leftmost index of any occurrence of u or v in E .
2. Let j be the rightmost index of any occurrence of u or v in E .
3. The LCA is $E[k]$, where $i \leq k \leq j$ and $E[k]$ is the node with the minimum *depth* in $E[i..j]$.

The first two steps can be done in $\mathcal{O}(1)$ after an $\mathcal{O}(n)$ preprocessing; we can simply store the leftmost and rightmost occurrence of every node in the Euler tour. The last step is a range minimum query on the *depth sequence* of the Euler tour.

Exercise 14. Prove that the above procedure correctly computes the $\text{lca}(u, v)$.

E%erci5e 14. Prove that the above procedure correctly compute5 the lca(u,v).

Let E be the Euler tour of the rooted tree and D the corresponding depth array. Let i and j be the leftmost and rightmost occurrences of u and v in E (assuming $i \leq j$).

Between i and j in the tour, the path goes up from u to their LCA and down to v , visiting the LCA and possibly other ancestors.

The LCA is the node with minimum depth among those in $E[i..j]$, since all nodes in this interval lie on or between the two root-to-leaf paths and the LCA is the highest common point.

Thus, $\text{lca}(u, v) = E[k]$ where k is the index of the minimum value in $D[i..j]$.

E%erci5e 15. Show that the length of E and D is $\Theta(n)$.

In an Euler tour of a tree with n nodes, each edge is traversed twice: once upward, once downward: A tree has $n - 1$ edges, so the tour contains $2(n - 1)$ edge traversals.

Adding the n nodes (each time a node is visited, it is recorded), the total length of E (and thus D), is:

$$n + 2(n - 1) = 3n - 2 = \Theta(n).$$