

1.1. Prove or Disprove: For every pair of nodes a and b (not necessarily distinct), node a is an ancestor of node b if and only if $\text{lca}(a,b) = a$. Note that a node a is considered an ancestor of itself.

[\Rightarrow] If $\text{lca}(a,b) = a$, then a is the lowest common ancestor of a and b, then in particular, node a is an ancestor of node b.

[\Leftarrow] If node a is an ancestor of node b, then a is a candidate for the lowest common ancestor, and since $\text{lca}(a,b)$ is the actual lowest ancestor, it is either a or something “lower” than a, but anything lower than a is not an ancestor of a. Therefore, $\text{lca}(a,b) = a$.

1.2. Prove or Disprove: For every triple of nodes a,b, and c, which are not necessarily distinct, the set $\{\text{lca}(a,b), \text{lca}(b,c), \text{lca}(c,a)\}$ has cardinality at most 2.

True. Let the three LCA values be sorted by depth such that $\text{depth}(X) \leq \text{depth}(Y) \leq \text{depth}(Z)$ for $\{X,Y,Z\} = \{\text{lca}(a,b), \text{lca}(b,c), \text{lca}(c,a)\}$.

Observe that Y and Z are both ancestors of c since each pair is the LCA of a pair involving c, so Y is an ancestor of Z.

Since Z is an ancestor of a, Y is also an ancestor of a. Thus, Y is a common ancestor of a and b. Since X is the LCA of a and b, we have $\text{depth}(X) \leq \text{depth}(Y)$, but with $\text{depth}(X) = \text{depth}(Y)$, and both being ancestors of b at the same depth, they must be the same node $X = Y$.

Hence, the set $\{\text{lca}(a,b), \text{lca}(b,c), \text{lca}(c,a)\}$ contains at most two distinct nodes. This equality holds if $\text{lca}(a,b) \neq \text{lca}(c,a)$, in which case the cardinality is 1.

Problem 1. There are n attractions in a city and m routes that connect these attractions. Each route connects two different attractions, is bidirectional, and takes a certain number of minutes to go through. There is exactly one way to go from any attraction to any other attraction via routes, if you’re not allowed to reuse a route.

You want to plan a tour around this city, but you want to be as time-efficient as possible. Answer q questions of the following form:

“What is the shortest amount of time it will take to go from attraction A to attraction B?

Note that the problem is *online*, you must answer each query without knowing the future queries.

2.1. Find an algorithm that solves Problem 1 where preprocessing takes $O(1)$ time, and each query takes $O(n)$ time.

We observe that we can model the problem as a graph, a tree with unique paths.

For each query(A,B), we run:

1. BFS/DFS from A until reaching B, tracking the unique path.
2. Summing the edge weights along this path using DP.

...

```
function query(A,B):  
    visited = array of size n, initialized to false  
    parent = array of size n, initialized to -1  
    cost_to = array of size n, initialized to 0  
  
    queue = [A]  
    visited[A] = true  
  
    while queue not empty:  
        u = queue.pop()  
        if u == B:  
            return cost_to[B]  
        for each neighbor v of u with weight w:  
            if not visited[v]:  
                visited[v] = true  
                parent[v] = u  
                cost_to[v] = cost_to[u] + w  
                queue.push(v)  
    ...
```

The time and space complexity is $O(n)$ because running DFS/BFS costs $O(n)$ and it takes $O(n)$ to initialize the required arrays.

2.1. Find an algorithm that solves Problem 1 where preprocessing takes $O(n^2)$ time, and each query takes $O(1)$ time.

For each node u:

- Run BFS/DFS from u to compute the shortest path distance to every other node.
- Use DP: $\text{dist}[u][v] = \text{dist}[u][\text{parent}] + w(u, \text{parent})$
- Store all distances in an $n \times n$ table.

```Pseudocode:

preprocess():

```
 dist = 2D array of size n x n
 for each node u in V:
 run BFS from u:
 dist[u][u] = 0
 for each node v visited:
 dist[u][v] = dist[u][parent[v]] + weight(parent[v], v)
 return dist
```

query(A,B,dist):

```
 return dist[A][B]
```

...

Preprocessing takes  $n * O(n) = O(n^2)$ , with table lookups being  $O(1)$ .

**2.3. Find an algorithm that solves Problem 1 where preprocessing takes  $O(n \log n)$  time, and each query takes  $O(\log n)$  time.**

We execute the following preprocessing:

1. Root the tree arbitrarily.
2. Compute depth and parent for each node via DFS:  $O(n)$
3. Precompute binary lifting table  $\text{up}[u][k]$ :  $2^k$ th ancestor of u:  $O(n \log n)$
4. Compute distance from root:  $\text{dist}[u] = \text{dist}[\text{parent}] + w(\text{parent}, u)$ :  $O(n)$

Query(A,B):

1.  $L = \text{LCA}(A, B)$  using binary lifting:  $O(\log n)$
2. Return  $\text{dist}[A] + \text{dist}[B] - 2 * \text{dist}[L]$

```Pseudocode:

preprocess():

```
DFS(root): compute parent, depth, dist[root] = -  
for k = 1 to log2(n)  
for each node u:  
    up[u][k] = up[up[u][k-1]][k-1]
```

lca(u,v):

```
if depth[u] < depth[v]: swap  
raise u to depth[v] using binary lifting  
if u == v: return u  
for k = log2(n) down to 0:  
    if up[u][k] != up[v][k]:  
        u = up[u][k]; v = up[v][k]  
return up[u][0]
```

query(A,B):

```
L = lca(A,B)  
return dist[A] + dist[B] - 2 * dist[L]
```

Therefore, $\langle O(n \log n), (\log n) \rangle$.

Problem 2. There are n attractions in a city and m routes that connect these attractions. Each route connects two different attractions, is bidirectional, and takes a certain number of minutes to go through. There is exactly one way to go from any attraction to any other attraction via routes, if you're not allowed to reuse a route.

You want to plan a tour around this city via electric car, but its battery only lasts for a limited number of minutes. There is a charging station at each attraction, where you can fully charge up your electric car's battery. Answer q questions of the following form:

“If you want to go from attraction A to attraction B, at least how many minutes should your battery be able to last?”

Note that the problem is online; you must answer each query without knowing the future queries.

3.1. Find an algorithm that solves Problem 2 that runs in $<O(1), O(n)>$.

First, we observe that the graph is a tree, and we can model it as an undirected weighted graph with a node for each attraction and an edge for each route, with cost equal to the amount of time to go through that route.

We do not preprocess anything, resulting in $O(1)$ time complexity.

For query (A,B) , we need the maximum edge weight along the unique path from A to B .

1. Run BFS/DFS from A until reaching B .
2. Track the maximum edge weight along the path via DP: $\text{max_to}[v] = \max(\text{max_to}[\text{parent}], w(\text{parent}, v))$

```Pseudocode:

```
function query(A,B):
 visited = array of size n, initialized to false
 max_to = array of size n, initialized to 0

 queue = [A]
 visited[A] = true

 while queue not empty:
 u = queue.pop()
 if u == B:
 return max_to[B]
 for each neighbor v of u with weight w:
 if not visited[v]:
 visited[v] = true
 max_to[v] = max(max_to[u], w)
 queue.push(v)
 ...
```

$O(n)$  per query.

### 3.2. Find an algorithm that solves Problem 2 that runs in $\langle O(n^2), O(1) \rangle$ .

This time, we preprocess for each node in  $V$ :

- Run BFS/DFS for  $u$  to compute the maximum edge weight along the path to every other node.
- Use DP:  $\text{max\_edge}[u][v] = \max(\text{max\_edge}[u][\text{parent}], w(\text{parent}, v))$
- Store all values in an  $n \times n$  table.

...

preprocess():

```
max_edge = 2D array of size n x n
for each node u in V:
 run BFS from u:
 max_edge[u][u] = 0
 for each node v visited:
 p = parent[v]
 max_edge[u][v] = max(max_edge[u][p], weight(p,v))
return max_edge
```

query(A,B,max\_edge):

```
return max_edge[A][B]
```

Preprocessing takes  $n * O(n) = O(n^2)$ , and queries take  $O(1)$  table lookup.

### 3.3. Find an algorithm that solves Problem 2 that runs in $\langle O(n\log n), O(\log n) \rangle$

Preprocessing goes as follows:

1. Root the tree arbitrarily. Run DFS to compute depth and parent.
2. Compute binary lifting tables:  $O(n\log n)$ 
  - a.  $\text{up}[u][k]$ :  $2^k$ -th ancestor of  $u$
  - b.  $\text{max}[u][k]$ : maximum edge weight from  $u$  up to its  $2^k$ -th ancestor
3. Recurrence:
  - a.  $\text{max}[u][0] = w(u, \text{parent})$
  - b.  $\text{max}[u][k] = \max(\text{max}[u][k-1], \text{max}[\text{up}[u][k-1]][k-1])$

Query(A,B):

1. Compute  $L = \text{LCA}(A, B)$  using binary lifting.
2. Compute maximum edge weight on path  $A \rightarrow L$ :
  - a. Climb from  $A$  to  $L$ , tracking max using max table
3. Compute maximum on path  $B \rightarrow L$  similarly as 2

4. Return the maximum of both sides

```Pseudocode:

```
max_on_path(u): // v is ancestor of u
    res = 0
    diff = depth[u] - depth[v]
    for k = log2(n) down to 0:
        if diff >= 2k:
            res = max(res, max[u][k])
            u = up[u][k]
            diff -= 2k
    return res

query(A,B):
    L = lca(A,B)
    return max(max_on_path(A,L), max_on_path(B,L))
````
```

Preprocessing: O( $n \log n$ )

Querying: O( $\log n$ )